

LECTURE 7

QUERY PROCESSING

Muhammad Hamiz Mohd Radzi
Faiqah Hafidzah Halim



Contents

- Overview of Query Processing
- Query Decomposition

Objectives

- At the end of this lesson, you should be able to:
 - *Define Query Processing, Query Optimization and Query Decomposition.*
 - *Compare different strategies of Query Processing.*
 - *Explain phases of Query Processing.*
 - *Explain advantages and disadvantages of dynamic and static optimization.*
 - *Describe the stages involve in Query Decomposition.*

Introduction

- When the relational model was first launched commercially, one of the major criticisms often cited was **inadequate performance of queries**.
- Since then, a significant amount of research has been devoted to **developing highly efficient algorithms** for processing queries.
- There are many ways in which a complex query can be performed, and one of **the aims** of query processing is to **determine which one is the most cost effective**.

- It is **programmer's responsibility** to select the most appropriate execution strategy.
- SQL (declarative language) user specifies what data is required rather than how it is to be retrieved. (to make SQL more universally usable)
- Giving DBMS the responsibility for selecting best strategy and more control over system performance.

Overview of Query Processing

- QP : Activities involved in **parsing, validating, optimizing and executing** a query.
- Aim of QP: To **transform** a query written in a **high level language** (SQL) into a **correct and efficient execution strategy** expressed in a **low-level language** (implementing Relational Algebra) and to **execute** the strategy to retrieve required data.
- An important aspect of query processing is **query optimization**.

Query Optimization (QO)

- QO: The activity of **choosing an efficient execution strategy** for processing a query.
- Generally, we try to reduce the **total execution time of the query**, which is the sum of the execution times of all individual operations that make up the query.
- However, **resource usage** may also be viewed as the **response time of the query**, in which case we concentrate on maximizing the number of parallel operations.

- Since the problem is **computationally hard to control** with a large number of relations, the strategy adopted is generally reduced to finding a **near optimum solution**.
- The first technique of QO uses **heuristic rules** that order the operations in a query.
- The other technique **compares different strategies** based on their relative costs and selects the one that minimizes resource usage.
- Both methods of query optimization depend on **database statistics** to evaluate properly the different options that are available.

- The **accuracy and currency** of these statistics have a significant bearing on the **efficiency** of the execution strategy chosen.
- The statistics cover information about relations, attributes, and indexes.
- There are a lot of methods in order to keep the statistics which are updated regularly (problematic), periodic update or user specified.

Find all Managers who work at a London branch.

We can write this query in SQL as:

```
SELECT *  
FROM Staff s, Branch b  
WHERE s.branchNo = b.branchNo AND  
      (s.position = 'Manager' AND b.city = 'London');
```

Three equivalent relational algebra queries corresponding to this SQL statement are:

- (1) $\sigma_{(\text{position}='Manager') \wedge (\text{city}='London') \wedge (\text{Staff.branchNo}=\text{Branch.branchNo})}(\text{Staff} \times \text{Branch})$
- (2) $\sigma_{(\text{position}='Manager') \wedge (\text{city}='London')}(\text{Staff} \bowtie_{\text{Staff.branchNo}=\text{Branch.branchNo}} \text{Branch})$
- (3) $(\sigma_{\text{position}='Manager'}(\text{Staff})) \bowtie_{\text{Staff.branchNo}=\text{Branch.branchNo}} (\sigma_{\text{city}='London'}(\text{Branch}))$

Calculation Result

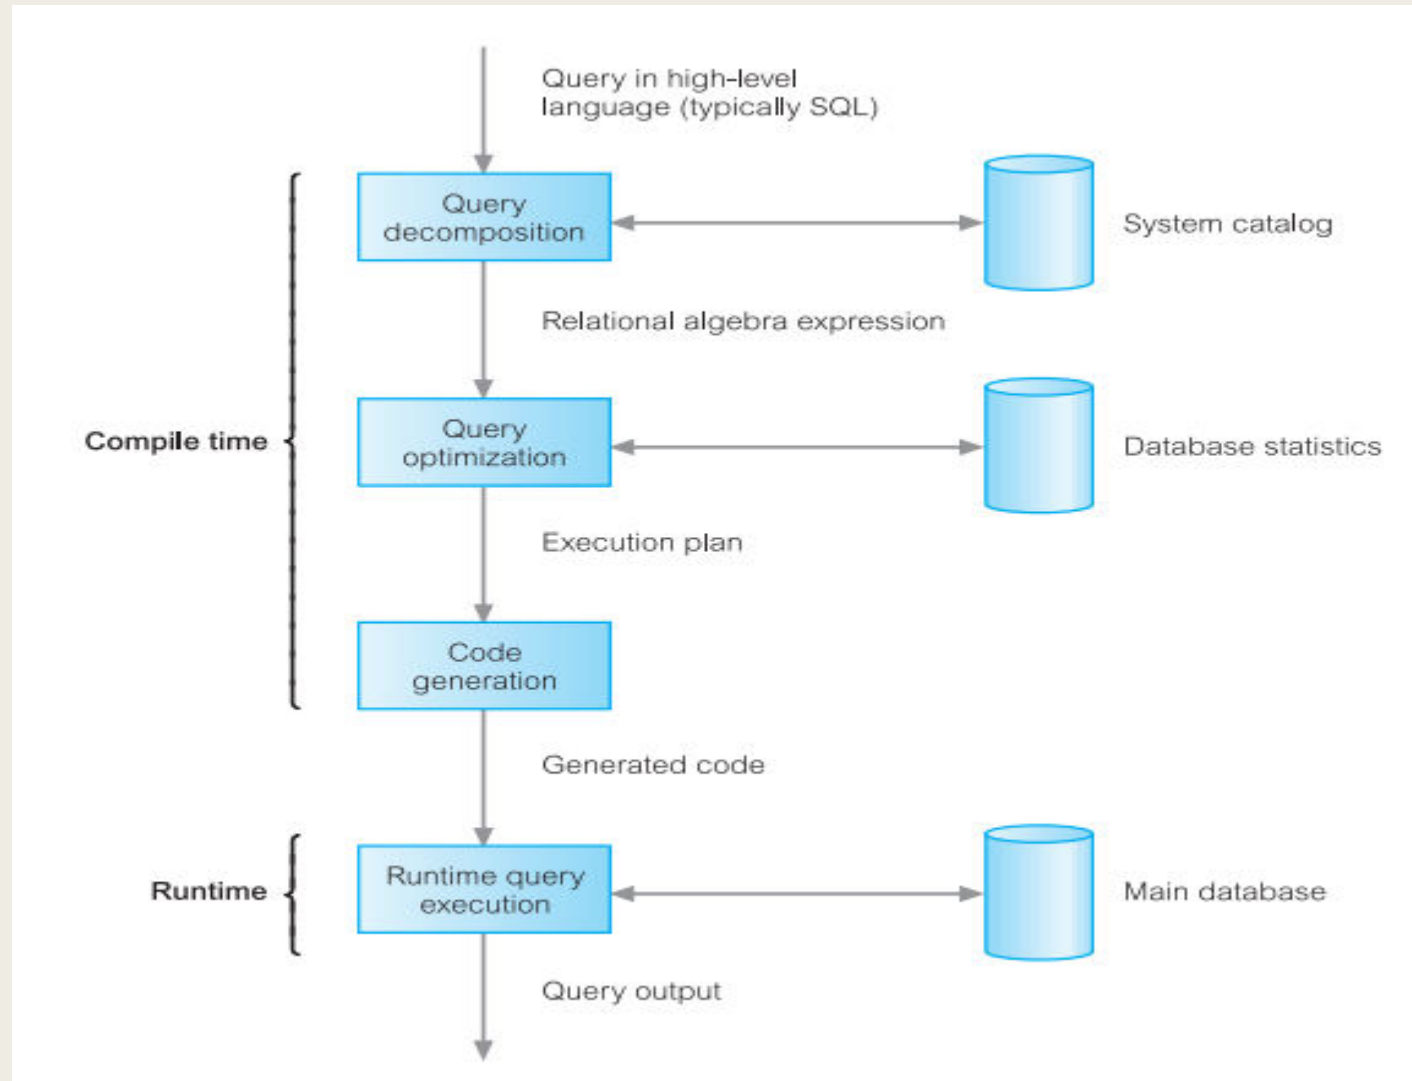
$$(1000 + 50) + 2*(1000 * 50) = 101\ 050 \text{ disk accesses}$$

- 1)

- $2*1000 + (1000 + 50) = 3050 \text{ disk accesses}$

- $1000 + 2*50 + 5 + (50 + 5) = 1160 \text{ disk accesses}$

Phases of Query Processing



Dynamic VS Static Processing

- There are two choices for when the first three phases of query processing can be carried out.
 1. **Dynamically** carry out decomposition and optimization every time the query is run.
 2. Just do it **once (statically)** for query to parsed, validated, and optimized

	ADVANTAGES	DISADVANTAGES
DYNAMIC	<p>All information required to select an optimum strategy is up to date.</p>	<p>Performance of the query is affected because the query has to be parsed, validated and optimized before it can be executed.</p> <p>May be necessary to reduce the number of execution strategies to be analyzed to achieve an acceptable overhead, which may have the effect of selecting a less optimum strategy.</p>
STATIC	<p>Runtime overhead is removed and more time available to evaluate a larger number of execution strategies (increasing the chances to find optimum strategy)</p> <p>For queries that are executed many times, taking some additional time to find a more optimum plan may prove to be highly beneficial.</p>	<p>Execution strategy is chosen to be optimal when the query is compiled may no longer be optimal when the query is running.</p>

Query Decomposition

- Query decomposition is the first phase of query processing.
- The **aims** of query decomposition are to **transform a high-level query into a relational algebra query**, and **to check that the query is syntactically and semantically correct**.
- The typical stages of query decomposition are
 - 1. Analysis**
 - 2. Normalization**
 - 3. Semantic Analysis**
 - 4. Simplification**
 - 5. Query restructuring**

Analysis


- The query is **lexically and syntactically analyzed** using the techniques of programming language compilers.
- Verifies that the **relations and attributes** specified in the query **are defined in the system catalog**.
- Verifies that any **operations applied** to database objects **are appropriate** for the object type.
- On completion of the analysis, the **high-level query** has been **transformed** into some **internal representation** (query tree) that is more suitable for processing.

Example

```
SELECT staffNumber  
FROM Staff  
WHERE position > 10;
```

This query would be rejected on two grounds:

- (1) In the select list, the attribute `staffNumber` is not defined for the `Staff` relation (should be `staffNo`).
- (2) In the `WHERE` clause, the comparison `'>10'` is incompatible with the data type `position`, which is a variable character string.



Normalization

- To transform the query into normalized form to facilitate further processing.
- Transformation rules are:
 - *Conjunctive NF: A sequence of conjuncts that are **connected with the \wedge (AND) operator**. Each conjunct contains one or more terms connected by the \vee (OR) operator.*
 - *Disjunctive NF: A sequence of disjuncts that are connected with **the \vee (OR) operator**. Each disjunct contains one or more terms connected by the \wedge (AND) operator.*

- Conjunctive normal form

$$(p_{11} \vee p_{12} \vee \cdots \vee p_{1n}) \wedge \cdots \wedge (p_{m1} \vee p_{m2} \vee \cdots \vee p_{mn})$$

- Disjunctive normal form

$$(p_{11} \wedge p_{12} \wedge \cdots \wedge p_{1n}) \vee \cdots \vee (p_{m1} \wedge p_{m2} \wedge \cdots \wedge p_{mn})$$

- Example: Find an employees who have been working for project P1 for 12 or 24 months.

■ SQL :

```
SELECT e.empName  
FROM emp e, assignment a  
WHERE e.empNo = s.empNo  
AND a.projNo = 'P1'  
AND duration = 12 OR duration = 24;
```

- Conjunctive:

$(e.\text{empNo} = a.\text{empNo}) \wedge (a.\text{projNo} = \text{'P1'}) \wedge (\text{duration} = 12 \vee \text{duration} = 24)$

- Disjunctive:

$(e.\text{empNo} = a.\text{empNo} \wedge a.\text{projNo} = \text{'P1'} \wedge \text{duration} = 12) \vee$

$(e.\text{empNo} = a.\text{empNo} \wedge a.\text{projNo} = \text{'P1'} \wedge \text{duration} = 24)$

Semantic Analysis

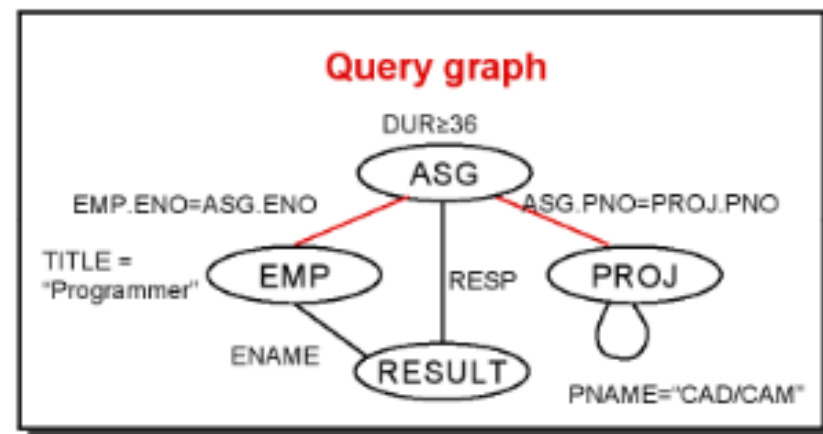
- Identify and reject **type incorrect** or **semantically incorrect** queries.
- Type incorrect: attributes and relations name **not in global schema** and **data type of attribute is not match**.
- Semantically incorrect: the relations are **not joined** in the **query and join graphs**.

- **Example:** Consider a query:

```

SELECT  ENAME, RESP
FROM    EMP, ASG, PROJ
WHERE   EMP.ENO = ASG.ENO
AND     ASG.PNO = PROJ.PNO
AND     PNAME = "CAD/CAM"
AND     DUR ≥ 36
AND     TITLE = "Programmer"

```

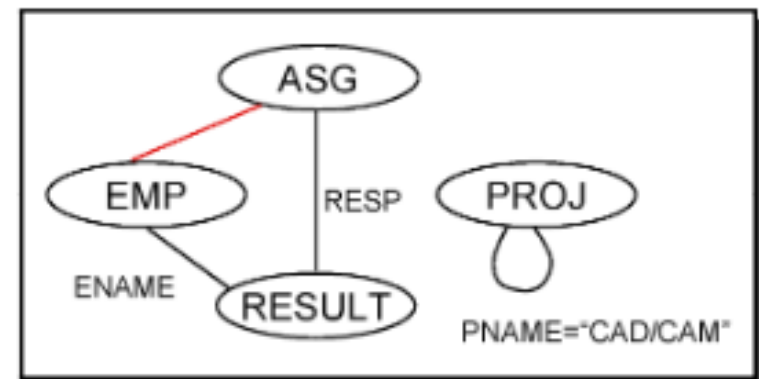


- **Example:** Consider the following query and its query graph:

```

SELECT  ENAME, RESP
FROM    EMP, ASG, PROJ
WHERE   EMP.ENO = ASG.ENO
AND     PNAME = "CAD/CAM"
AND     DUR ≥ 36
AND     TITLE = "Programmer"

```



Simplification

- To detect **redundant qualifications**, **eliminate common sub-expressions**, and **transform** the query to a **semantically equivalent** but more easily and efficiently computed form.
- Access restrictions, view definitions, and integrity constraints are considered at this stage.

Query Restructuring

- In the final stage of query decomposition, the query is restructured to provide a more efficient implementation.

Reference

- *Database Systems: A Practical Approach to Design, Implementation, and Management*, Thomas Connolly and Carolyn Begg, 5th Edition, 2010, Pearson.